



APACHE SLING & FRIENDS TECH MEETUP  
BERLIN, 25-27 SEPTEMBER 2017

# Building an Apache Sling Rendering Farm

Bertrand Delacretaz

@bdelacretaz

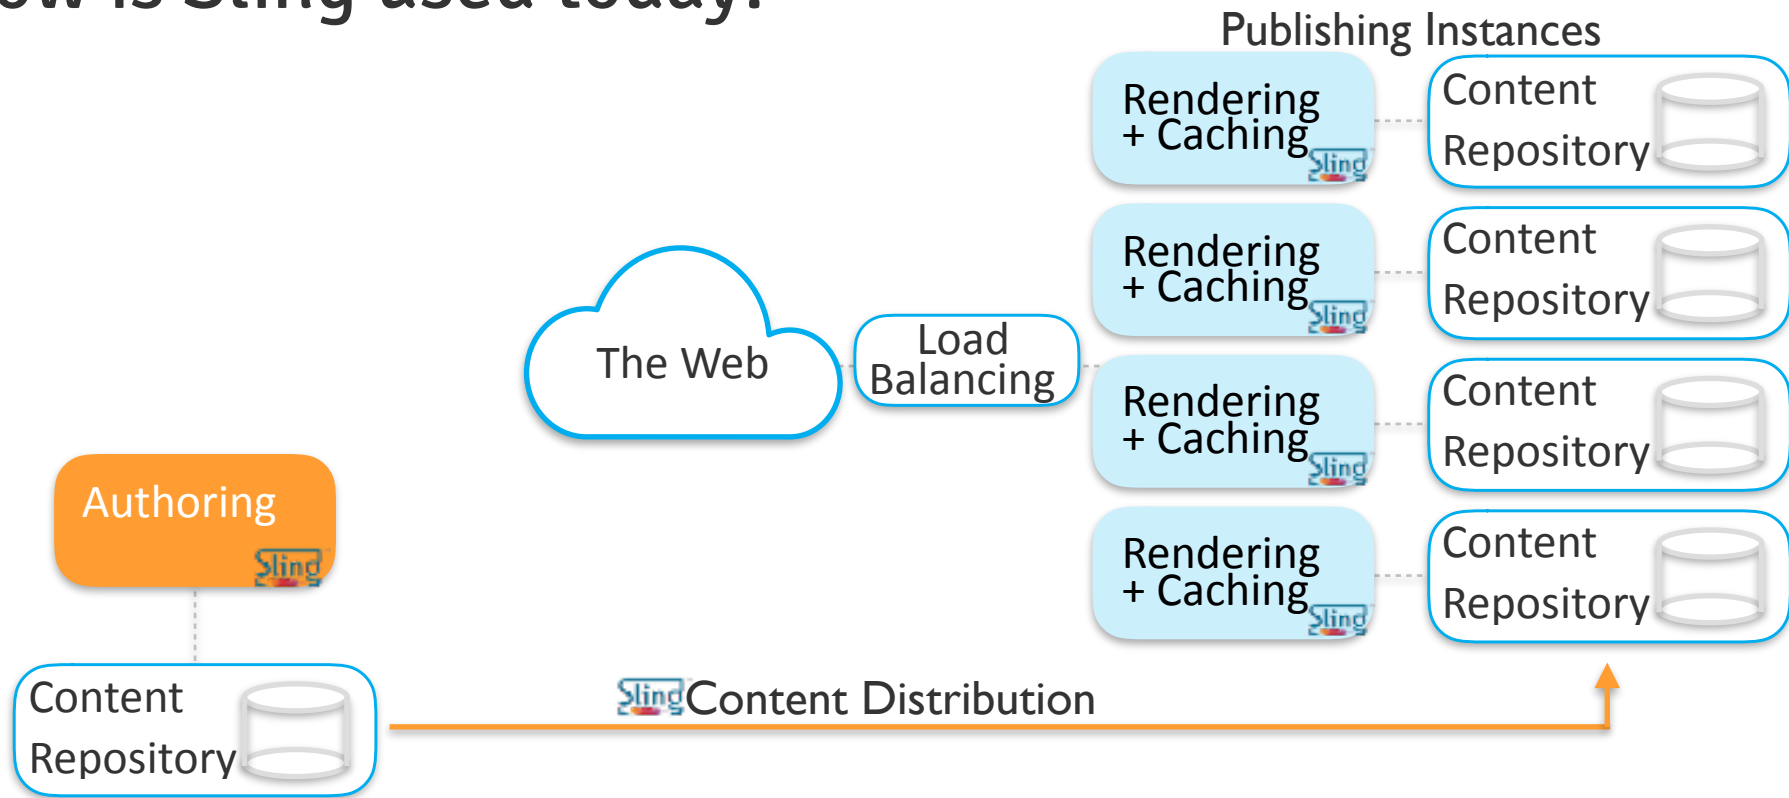
Sling committer and PMC member  
Principal Scientist, Adobe AEM team



# What are we building?

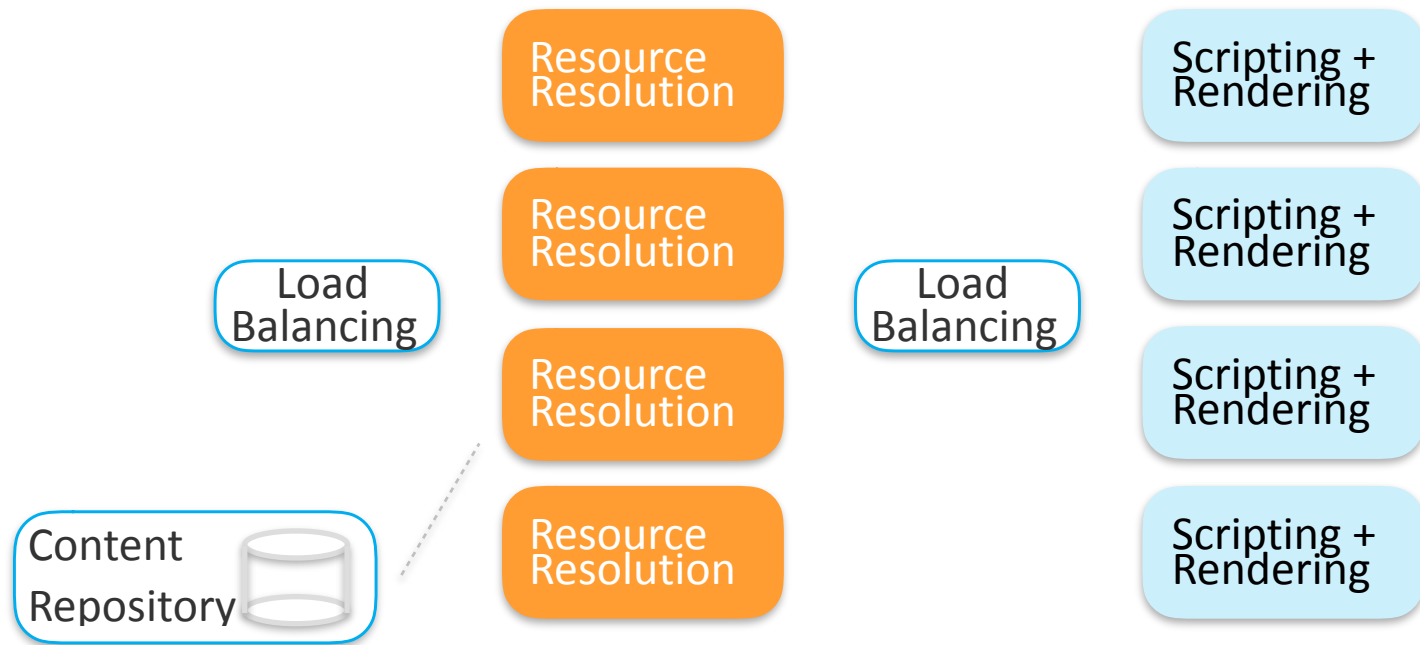
setting the stage

# How is Sling used today?



Sling instances dedicated to single tenants or “friendly” tenants.

# A Massive Sling Rendering/Processing Farm?

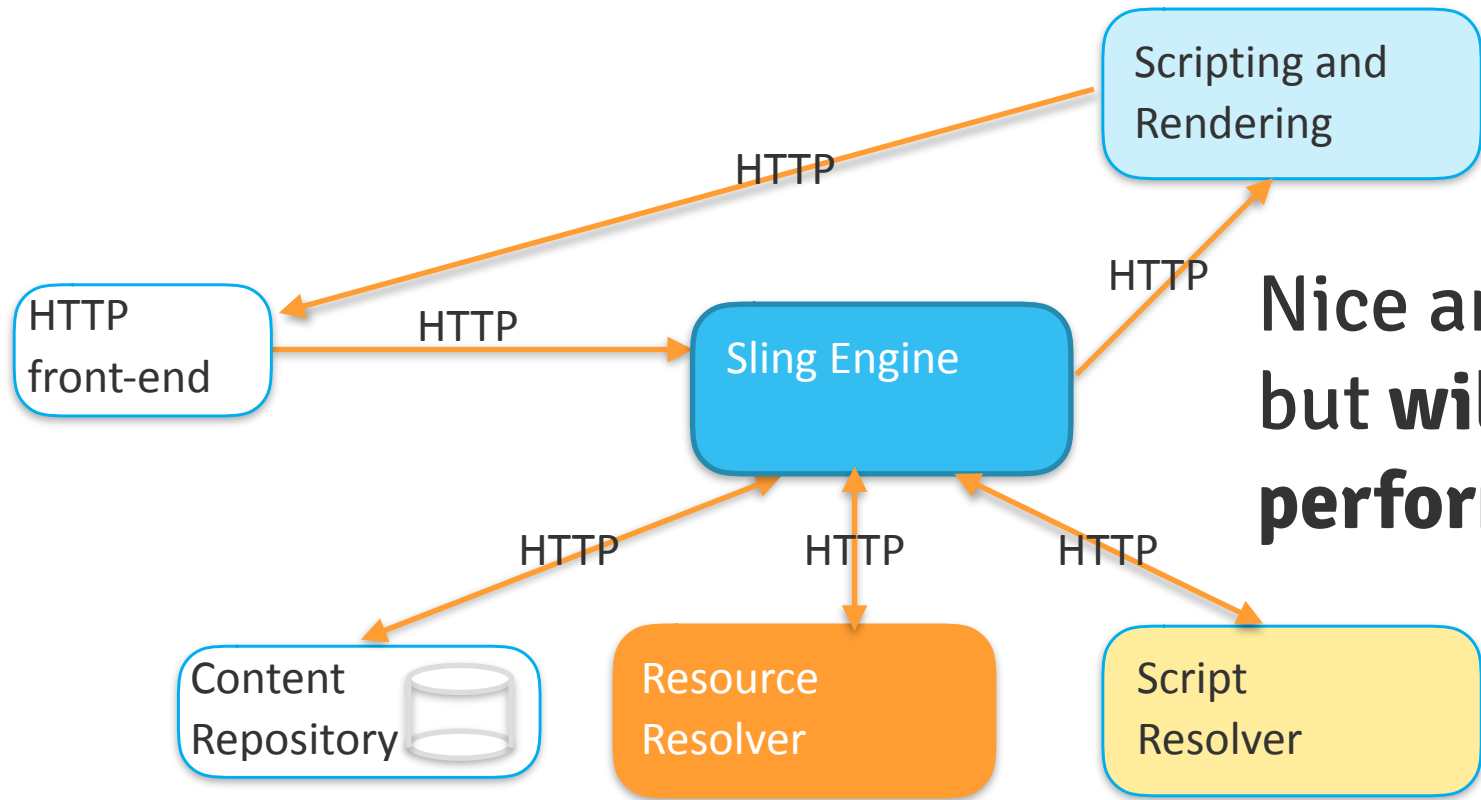


Elastic scaling at each stage  
Multiple developers (“tenants”) see their own world only

# Federated Services

This 2017 after all

# Microservices!



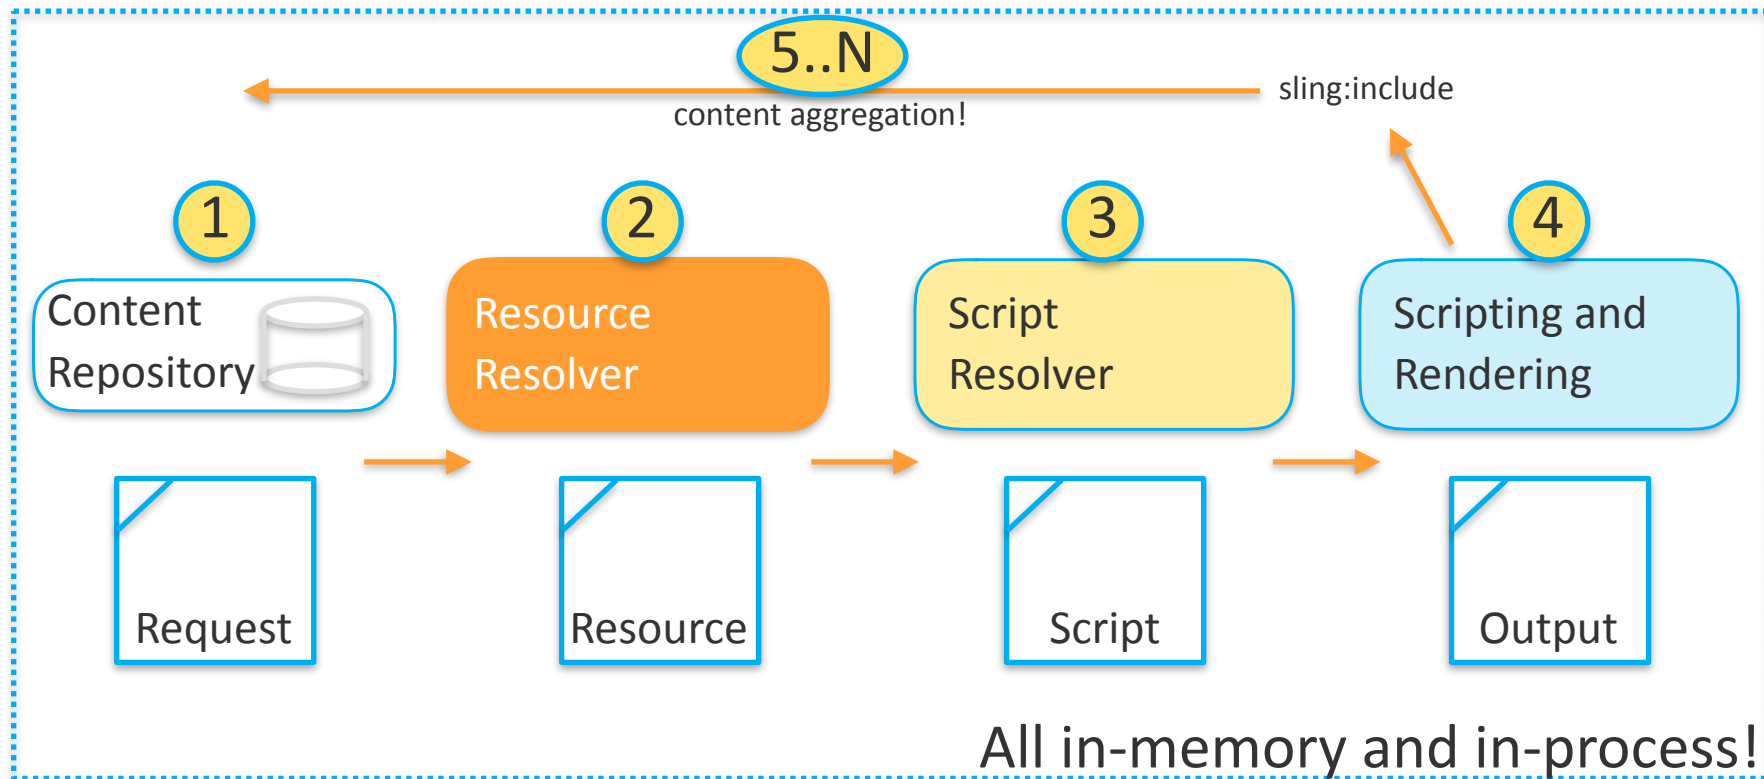
Nice and trendy,  
but **will that perform?**

Each component is an independent HTTP-based service, aka “religious microservices”

# The Sling Pipeline

Faithfully serving requests since 2007!

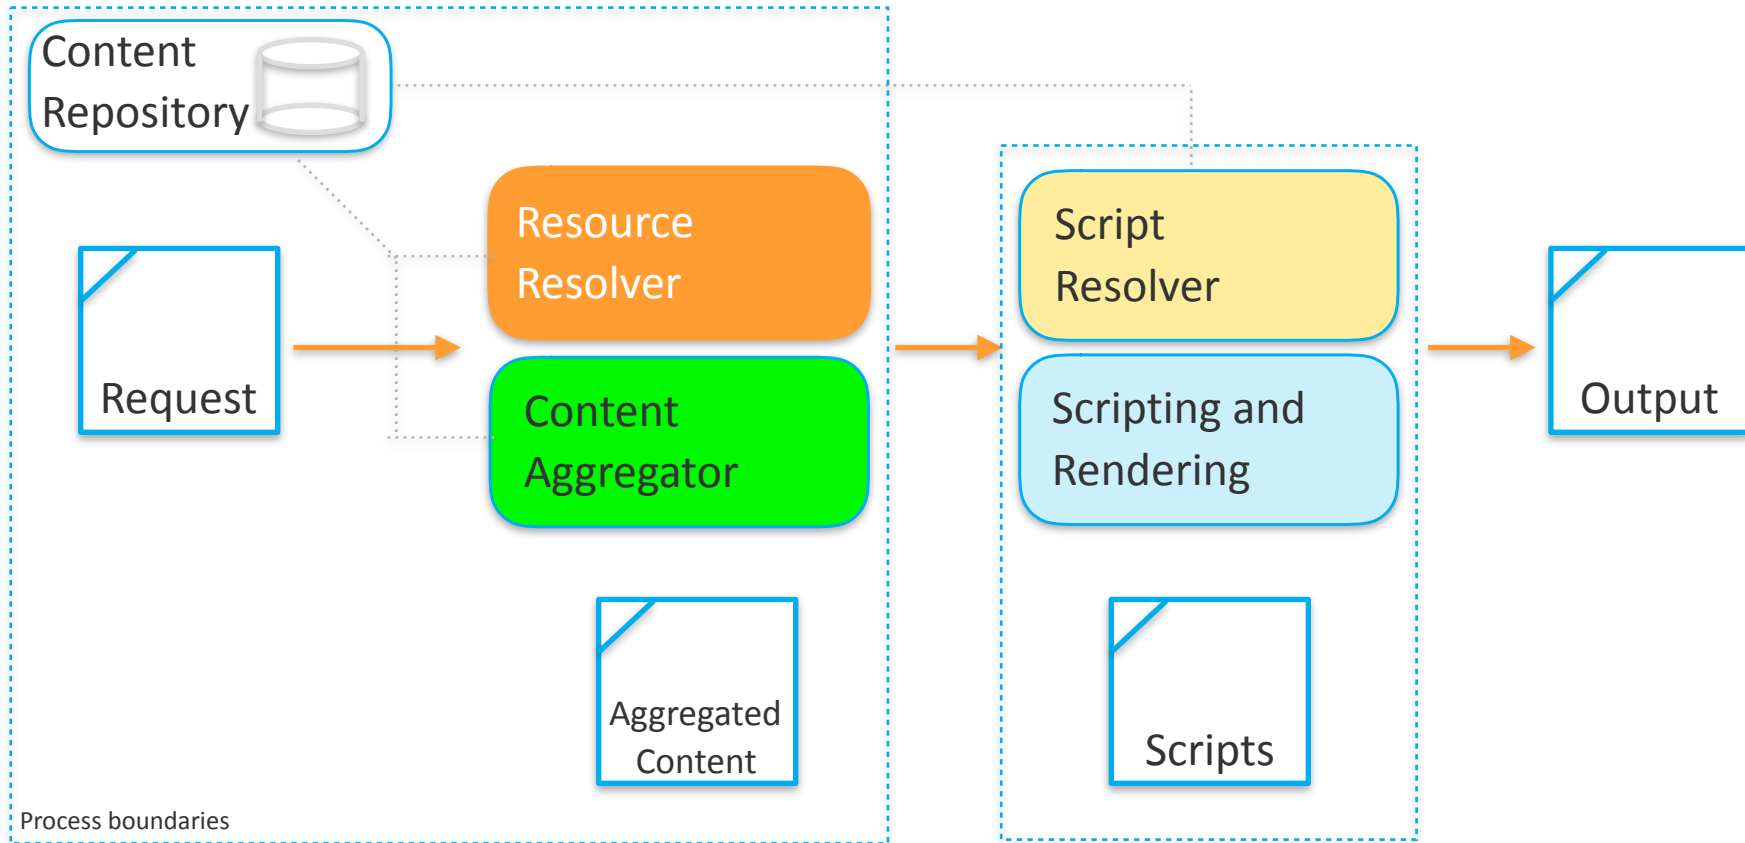
# Sling Request Processing Pipeline



Conceptually, the request hits the repository first, to get the Resource.  
Scripts and Servlets are equivalent, considering scripts only here.



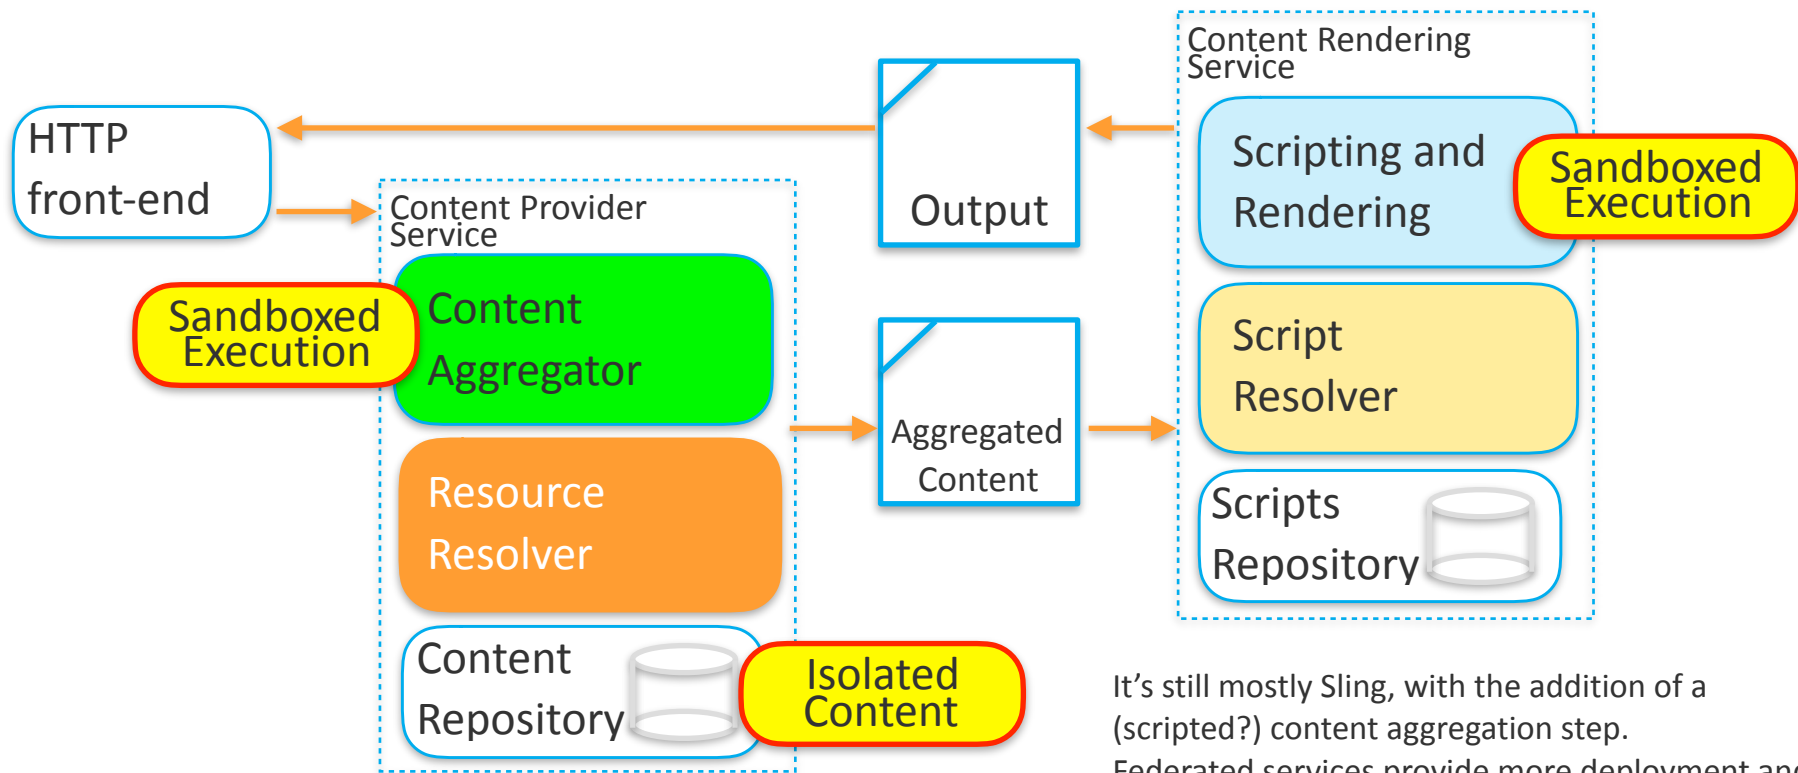
# Federated Services Friendly?



# Reasonably Federated?

Can we get isolation AND performance?

# Reasonably Federated Sling Rendering Farm?

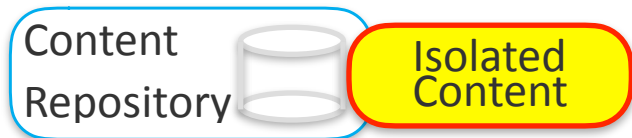


It's still mostly Sling, with the addition of a (scripted?) content aggregation step. Federated services provide more deployment and scaling options.

# Sandboxing & Isolation

How?

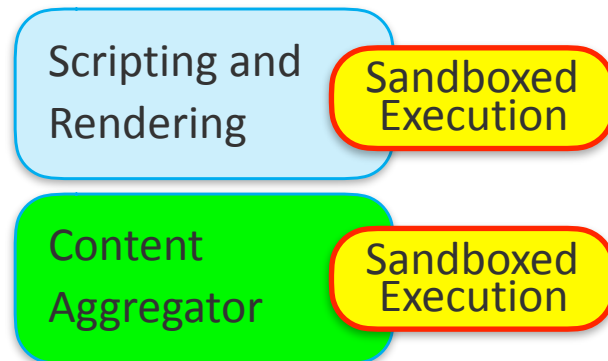
# Sandboxing & Isolation?



## Repository Access Control

can work but require a *dynamic search path* in Sling, see our experiments. Impacts caching, and mapping of incoming to resource paths is needed. Tried and tested.

**Repository jails** look possible with probable impact on Sling internals. Same with **multiple SlingRepository services**. New and more like a blacklist.

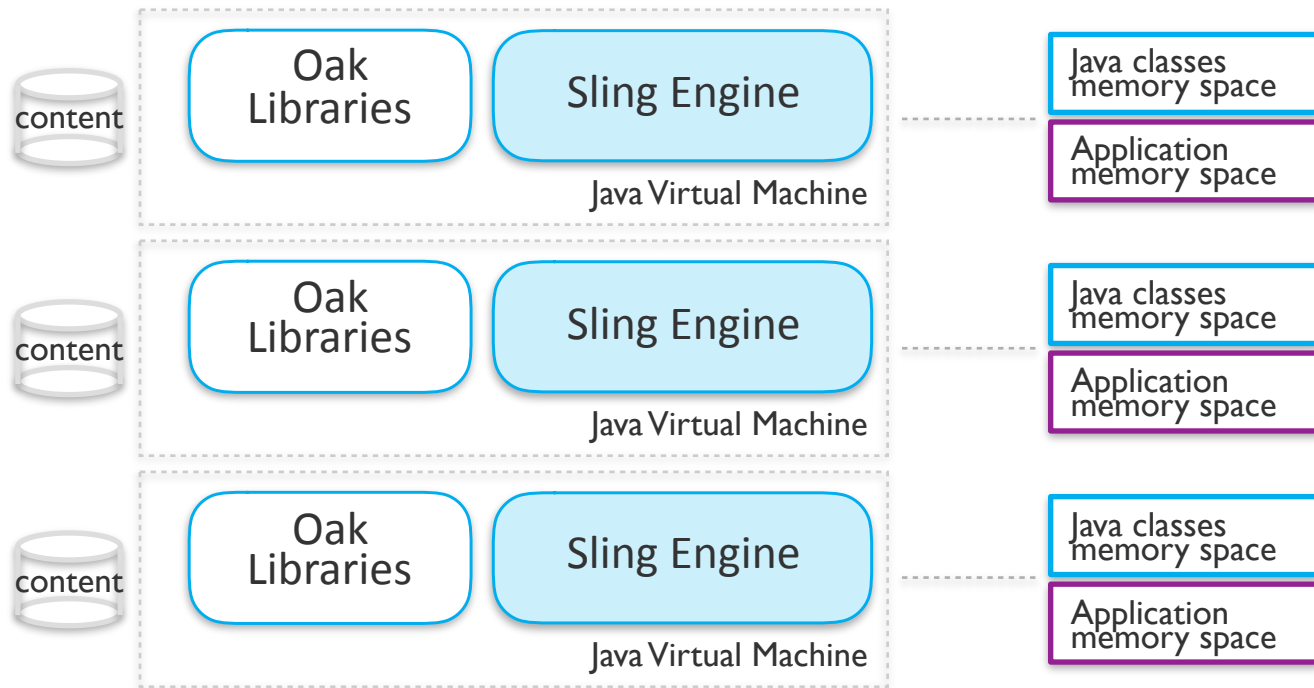


**Custom, restricted languages** are the safest? HTL (Use-API?), Handlebars?

**Sandboxing Nashorn** (JavaScript) looks possible but not ideal, see our experiments.

**Sandboxing Java** is *not realistic*- IBM canceled multi tenant JVM project for example.

# But it's a VM, right?



Perfect isolation!

But suboptimal use of resources!  
(and containers wouldn't help)

# Sandboxing scripting languages?

Java classes  
& services

OS  
Resources

```
<%  
var length = 0;  
if (request.getRequestParameter("file") != null) {  
    var file = null;  
    // store file  
    var reqPara = request.getRequestParameter("file");  
    var is = reqPara.getInputStream();  
    file = Packages.java.io.File.createTempFile("posttest", ".txt");  
    var fout = new Packages.java.io.FileOutputStream(file);  
    var c;  
    while ((c = is.read()) != -1) {  
        fout.write(c);  
    }  
    fout.close();  
  
    // read length  
    length = file.length();  
}  
%>
```

Infinite  
Loops

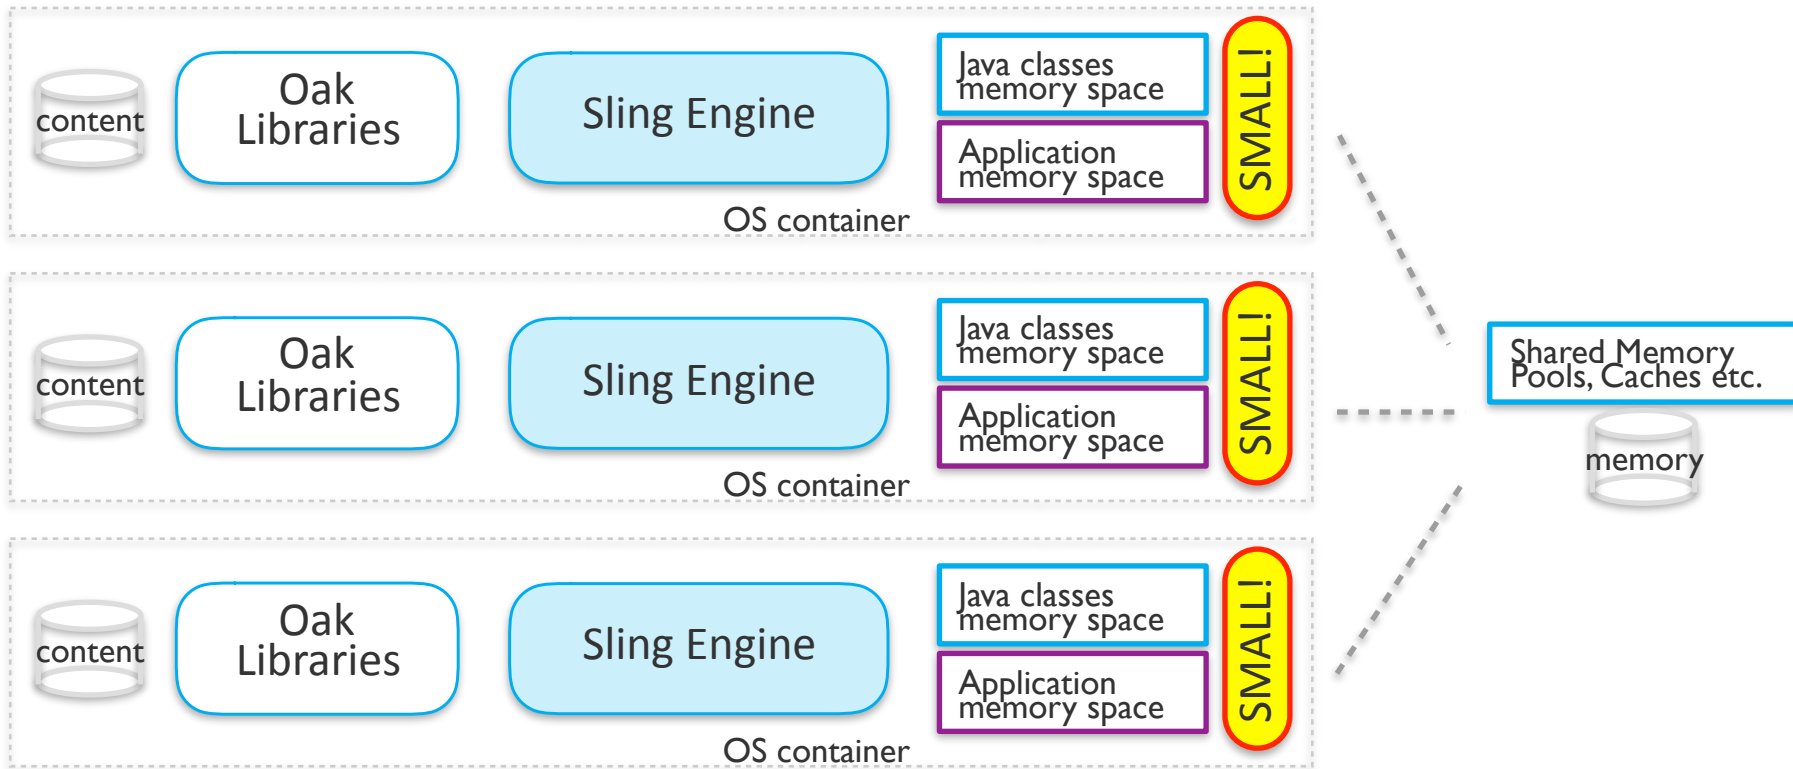
Memory  
Usage?

Many things need to be limited.

Whitelist approach is much safer -> custom languages?

HTL inherently sandboxed, except its Use-objects

# Containers?



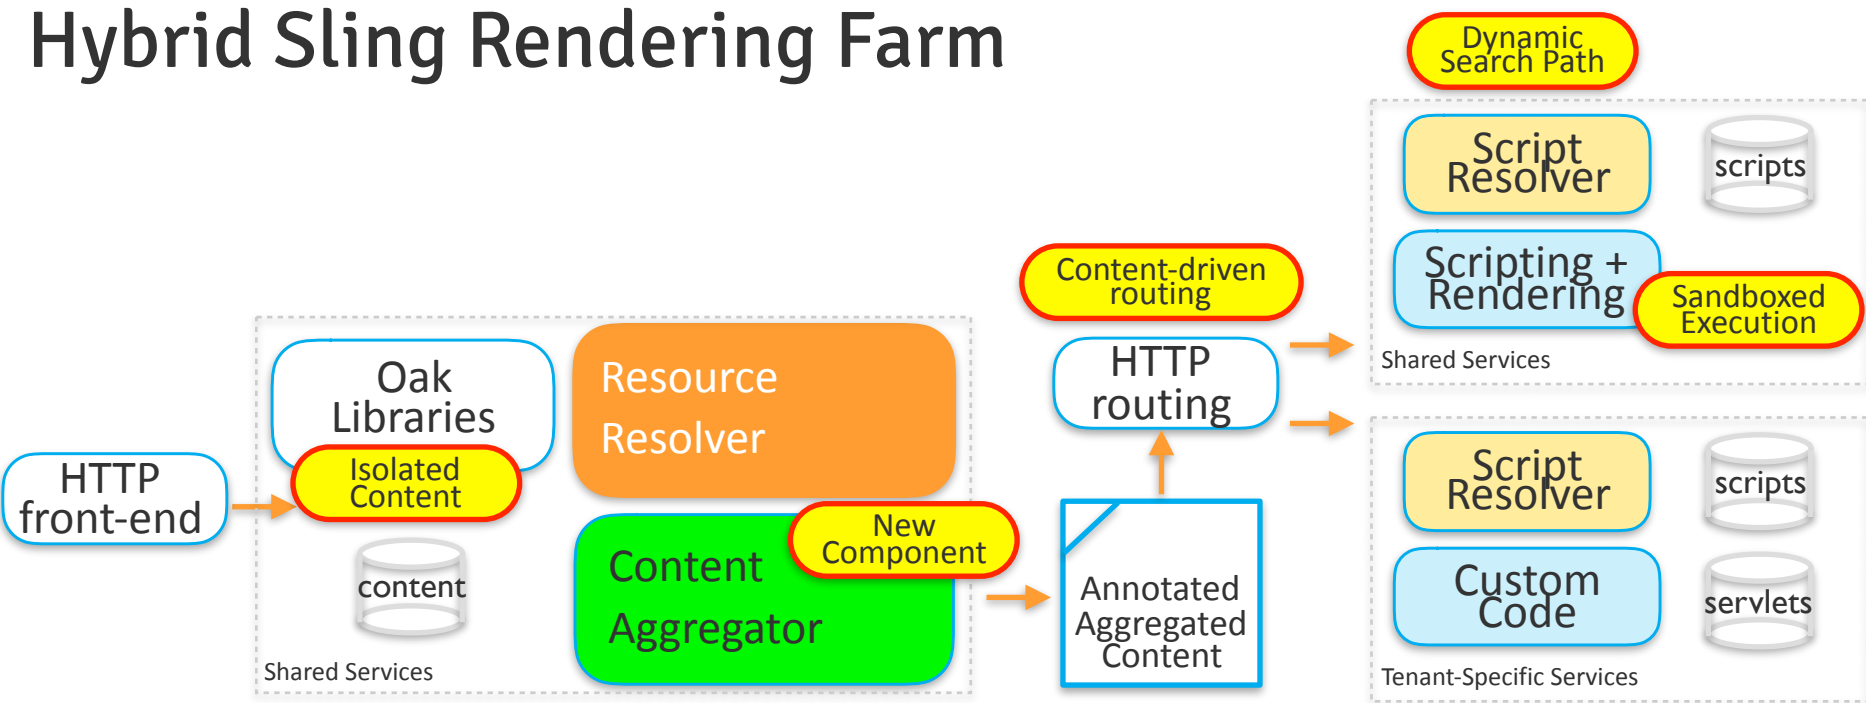
Same problem as multiple JVMs

Sharing caches, compiled scripts etc. can be a pragmatic solution.



# What do we do?

# Hybrid Sling Rendering Farm

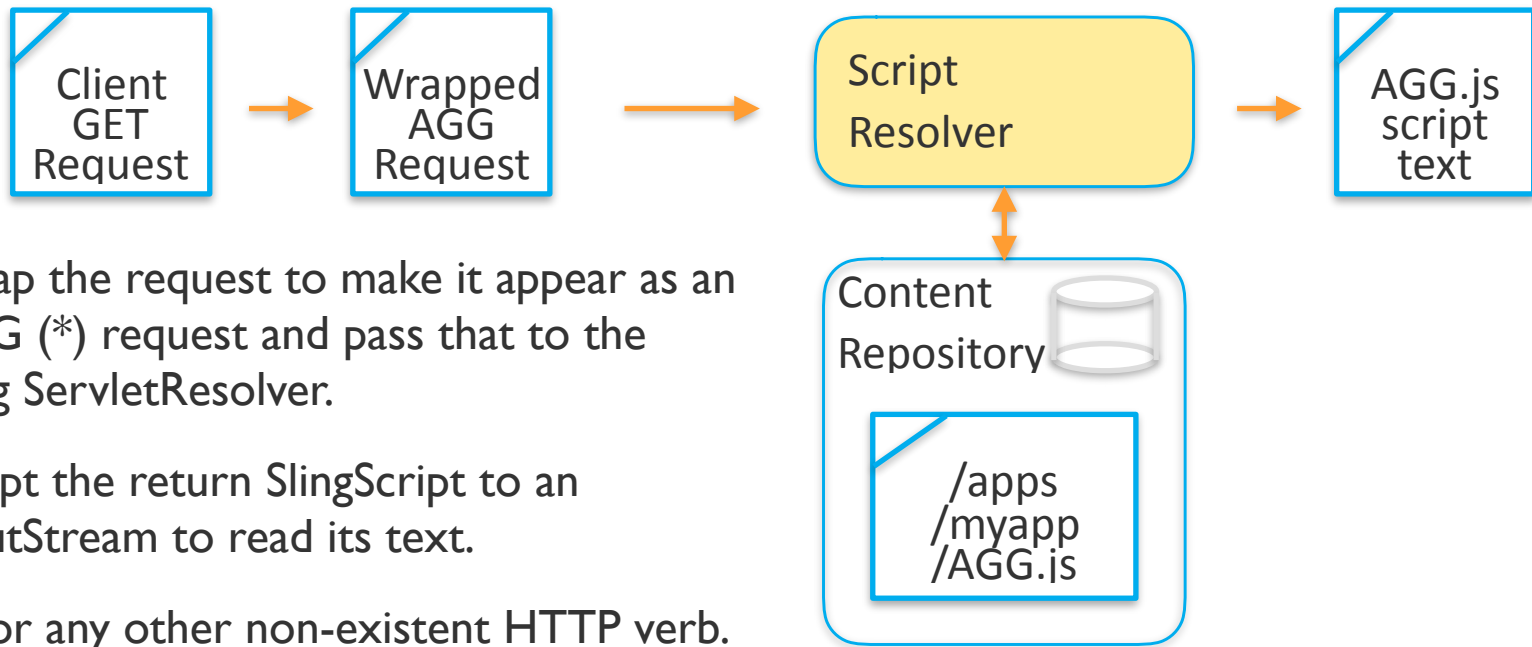


Provides the flexibility of Sling via tenant-specific services and dynamic routing.  
Uses shared services for the common parts.  
Allows for billable options depending on the actual routing.

# Experiments

building blocks that might be reusable

# Resolving new types of scripts



Wrap the request to make it appear as an AGG (\*) request and pass that to the Sling ServletResolver.

Adapt the return SlingScript to an InputStream to read its text.

(\*) or any other non-existent HTTP verb.

Code at <https://github.com/bdelacretaz/sling-adapto-2017> (ContentBVP.java)

# Resolving a SLING-CONTENT script

```
String getAggregatorScript(SlingHttpServletRequest r) {
    String result = null;
    Servlet s =
        servletResolver.resolveServlet(
            new ChangeMethodRequestWrapper(r, "SLING-CONTENT"));
    if(s instanceof SlingScript) {
        InputStream is = ((SlingScript)s).getScriptResource()
            .adaptTo(InputStream.class);    }
        if(is != null) {
            result = IOUtils.toString(is)
        }
    }
    return result;
}
```

Code at <https://github.com/bdelacretaz/sling-adappto-2017> (ContentBVP.java)

Experiment  
adaptTo() Bonus Points!

# Content Aggregation with Sling Query

```
var $ = Packages.org.apache.sling.query.SlingQuery.$
var SearchStrategy =
  Packages.org.apache.sling.query.api.SearchStrategy
var resourceResolver = resource.getResourceResolver()

var result = {
  siblings : $(resource).siblings(),

  rootChildren : $(resource).parents().last().children(),

  queryResult :
    $(resourceResolver)
      .searchStrategy(SearchStrategy.QUERY)
      .find("nt:base[title=foo]")
}
```

Content  
Aggregator

Sandboxed  
Execution

Used in a BindingsValuesProvider?

Or in a custom json renderer servlet which runs this script.

Inherently sandboxed due to custom language.

<https://sling.apache.org/documentation/bundles/sling-query.html>

# Dynamic scripts/servlet search path

```
if(dynamicServletResolver.canResolve(resource)) {  
    servlet = dynamicServletResolver.resolveServlet(request);  
} else {  
    ...existing resolver code  
}
```

A fairly simple change to the SlingServletResolver - should evolve into a real extension point if desired, and probably get the request as well.

Tested in SLING-4386 - another multitenant experiment which provides tenant-specific scripts but no real isolation.

Currently requires disabling the servlet resolution cache.

# Nashorn (JavaScript) sandboxing (Java Delight)

Experiment

```
NashornSandbox {  
    allow(final Class<?> clazz);  
    injectGlobalVariable(String variableName, Object object);  
    setMaxCPUTime(long limitMsec);  
  
    Object eval(final String javaScriptCode);  
  
    allowPrintFunctions(boolean v);  
    allowReadFunctions(boolean v);  
    ...more allow functions  
  
    // $ARG, $ENV, $EXEC...  
    allowGlobalsObjects(final boolean v);  
}
```

## Nashorn Sandbox

A secure sandbox for executing JavaScript in Java apps.

Also see [Rhino Sandbox](#).

Part of the [Java Delight Suite](#).

build passing

Uses Nashorn's **ClassFilter** to block Java classes

Sandboxing rewrites standard methods + user code- > **blacklisting**, not ideal

<https://github.com/javadelight/delight-nashorn-sandbox> (Java Delight Suite)



# CODA

where to now?

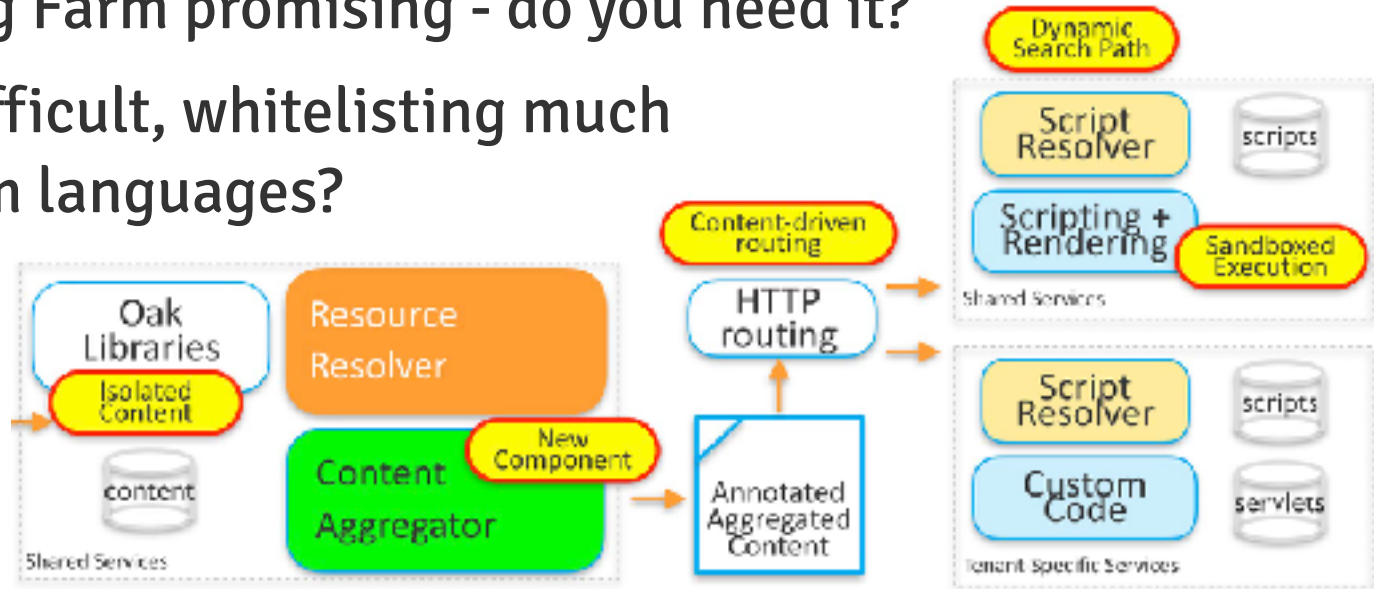
# CODA

In-memory nature of Sling is an important differentiator, in good and bad ways!

Hybrid Rendering Farm promising - do you need it?

Sandboxing is difficult, whitelisting much preferred, custom languages?

Reusable experiments?



Thank you for attending!

I'm Bertrand Delacretaz (@bdelacretaz)